


EN3160 L02

Basics and Point Operations

What is a Digital Image?

- A grayscale digital image is a rectangular array of numbers which represent pixels.
- Each pixel can take an integer value in the range $[0, 255]$, for an 8-bit image.
- If the image is a color image, then there would be three such arrays.
- The size of this array is actually the resolution of the image, e.g., example, 3712×5568 .
- We take the top-left pixel as the $(0, 0)$ pixel and vertical axis as the i axis.



What is a Color Digital Image?

- The grayscale image that we considered as a two-dimensional array, or a single plane.
- A color image has three such planes, one for blue, one for green and one for red. We call such an image an BGR image.
- If we access a pixel location such as (2, 3), we will get three values, B, G, and R.
- Each value is in $[0, 255]$. In this context, $2^8 \times 2^8 \times 2^8$ different colors are possible.




Image Resolution and DPI

- Image Resolution:
 - Number of pixels in an image (width × height)
 - Example: 1920×1080 (Full HD) = 2,073,600 pixels
 - Impacts the image detail and sharpness, file size and processing time
- Dots Per Inch (DPI):
 - Number of printed dots in one inch, i.e., physical point density
 - Typical values: screen: ~72–96 DPI, print: 300 DPI or higher

Creating a 6 × 8 Image


A Grayscale Image

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

im = np.zeros((6,8), dtype=np.uint8)

im[2,3] = 255

fig, ax = plt.subplots(1, 1, figsize=(6, 8))
ax.imshow(im, cmap='gray', vmin=0, vmax=255)
ax.xaxis.set_ticks_position('top')
plt.show()
```



A Color Image

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

im = np.zeros((6,8,3), dtype=np.uint8)

im[2,3] = (255, 190, 203)

fig, ax = plt.subplots(1, 1, figsize=(6, 8))
ax.imshow(im)
ax.xaxis.set_ticks_position('top')
plt.show()
```




Image Opening and Displaying

Displaying Using Matplotlib

```
import cv2 as cv
import matplotlib.pyplot as plt
im = cv.imread('images/jolie.png')
fig, ax = plt.subplots(1, 1, figsize=(6, 8))
ax.imshow(cv.cvtColor(im, cv.COLOR_BGR2RGB))
ax.xaxis.set_ticks_position('top')
plt.show()
```

Displaying Using OpenCV

```
import cv2 as cv
im = cv.imread('images/jolie.png')
cv.namedWindow('Image', cv.WINDOW_AUTOSIZE)
cv.imshow('Image', im)
cv.waitKey(0)
cv.destroyAllWindows()
```

Q: Why do we need to cv.cvtColor only when displaying using Matplotlib?

Displaying Image Properties

```
import cv2 as cv
import matplotlib.pyplot as plt
im = cv.imread('images/ryan.jpg')
fig, ax = plt.subplots(1, 1, figsize=(6, 8))
ax.imshow(cv.cvtColor(im, cv.COLOR_BGR2RGB))
ax.xaxis.set_ticks_position('top')
plt.show()
print('Image Shape:', im.shape)
print('Image Data Type:', im.dtype)
print('Image Size:', im.size)
```

Image Shape: (2641, 1761, 3)
Image Data Type: uint8
Image Size: 13952403

Increasing the Brightness

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

im1 = cv.imread('images/emma_gray.jpg',
cv.IMREAD_GRAYSCALE)

im2 = cv.add(im1, 100)

fig, ax = plt.subplots(1, 2, figsize=(6, 8))
ax[0].imshow(im1, cmap='gray', vmin=0, vmax=255)
ax[0].set_title('Original Image')
ax[1].imshow(im2, cmap='gray', vmin=0, vmax=255)
ax[1].set_title('Brightness Increased')
for a in ax:
    a.axis('off')
plt.show()
```

Original Image



Brightness Increased



What is wrong with this?

```
im2 = im1 + 100
```

Original Image



Brightness Increased



Increasing the Brightness Using Loops (Slow)

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

im1 = cv.imread('images/emma_gray.jpg', cv.IMREAD_GRAYSCALE)
im2 = np.zeros_like(im1)

for i in range(im1.shape[0]):
    for j in range(im1.shape[1]):
        im2[i,j] = im1[i,j] + 100

fig, ax = plt.subplots(1, 2, figsize=(6, 8))
ax[0].imshow(im1, cmap='gray', vmin=0, vmax=255)
ax[0].set_title('Original Image')
ax[1].imshow(im2, cmap='gray', vmin=0, vmax=255)
ax[1].set_title('Brightness Increased')

for a in ax:
    a.axis('off')

plt.show()
```

Obtaining One Color Plane

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt


im = cv.imread('images/rgb_flowers.jpg')
im_blue = im.copy()
im_blue[:, :, 1] = 0
im_blue[:, :, 2] = 0

fig, ax = plt.subplots(1, 2, figsize=(12, 8))
ax[0].imshow(cv.cvtColor(im, cv.COLOR_BGR2RGB))
ax[0].set_title('Original Image')
ax[1].imshow(cv.cvtColor(im_blue, cv.COLOR_BGR2RGB))
ax[1].set_title('Blue Channel')
for a in ax:
    a.axis('off')

plt.show()
```





Working of a Camera



1. Viewfinder (eyepiece lens)
2. Pentaprism
3. Focusing screen
4. Condenser lens
5. Color and infrared filter
6. Digital sensor
7. Shutter
8. Display
9. Electronics
10. Autofocus system
11. Reflex and relay mirror
12. Focusing elements
13. Aperture
14. Zoom elements
15. Front line gathering elements
16. Batteries


Bayer
arrangement of
color filters




Demosaicing:
Estimation of missing
components from
neighboring values




Color




Wavelengths comprising the visible range of the electromagnetic spectrum.



Color spectrum seen by passing white light through a prism.



Additive mixing







Subtractive mixing

A **color model** (color space or color system) facilitates the specification of colors: (1) a coordinate system, and (2) a subspace within that system, such that each color in the model is represented by a single point contained in that subspace.

RGB: image capturing in a camera, wavelength-based

CMYK (Cyan, Magenta, Yellow, Black): for printing

HSV (Hue, Saturation, Value): how humans describe color, decouples the color and gray-scale information




Summary

1. A grayscale digital image is an array of 8-bit unsigned integers in $[0, 255]$. We call each integer a pixel.
2. Color images have three such arrays (planes), one for red, one for green, and one for blue.
3. We can manipulate an image using OpenCV function or a pair of for-loops to access each pixel (slower).

Sampling and Reconstruction


Sampling and Reconstruction

- Simple example: a sine wave
- What if we “missed” things between the samples?
 - Unsurprising result: information is lost
 - Surprising result: indistinguishable from lower frequencies (or even higher frequencies)



Sampling and Reconstruction

- Simple example: a sine wave
- What if we “missed” things between the samples?
 - Unsurprising result: information is lost
 - Surprising result: indistinguishable from lower frequencies (or even higher frequencies)
 - *Aliasing*: signal “traveling in disguise” as other frequencies



Wagon Wheel Effect



<https://www.youtube.com/watch?v=6XwgbHjRo30>

Aliasing in Images


256x256

128x128

64x64



32x32

16x16



Aliasing “in the Wild”

Disintegrating textures



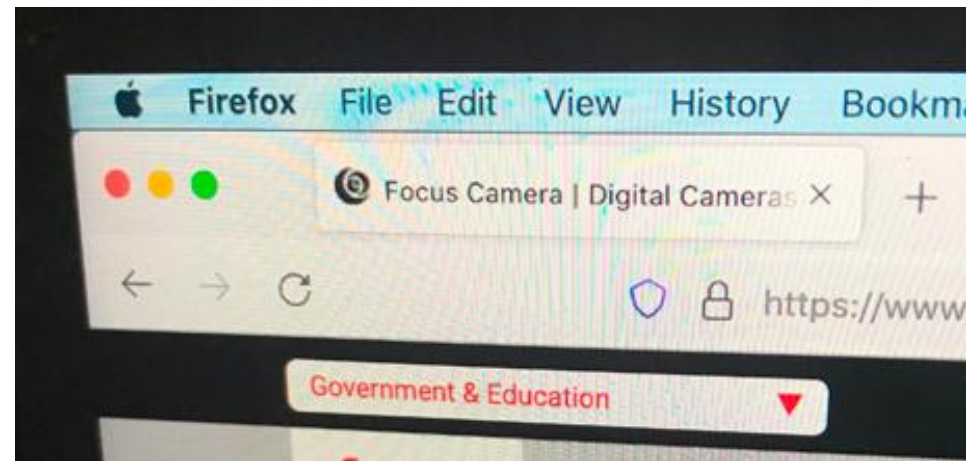
[Source](#)

Source: Lazeknik

Moire patterns, false color



[Source](#)



[Source](#)

Aliasing in Neural Networks

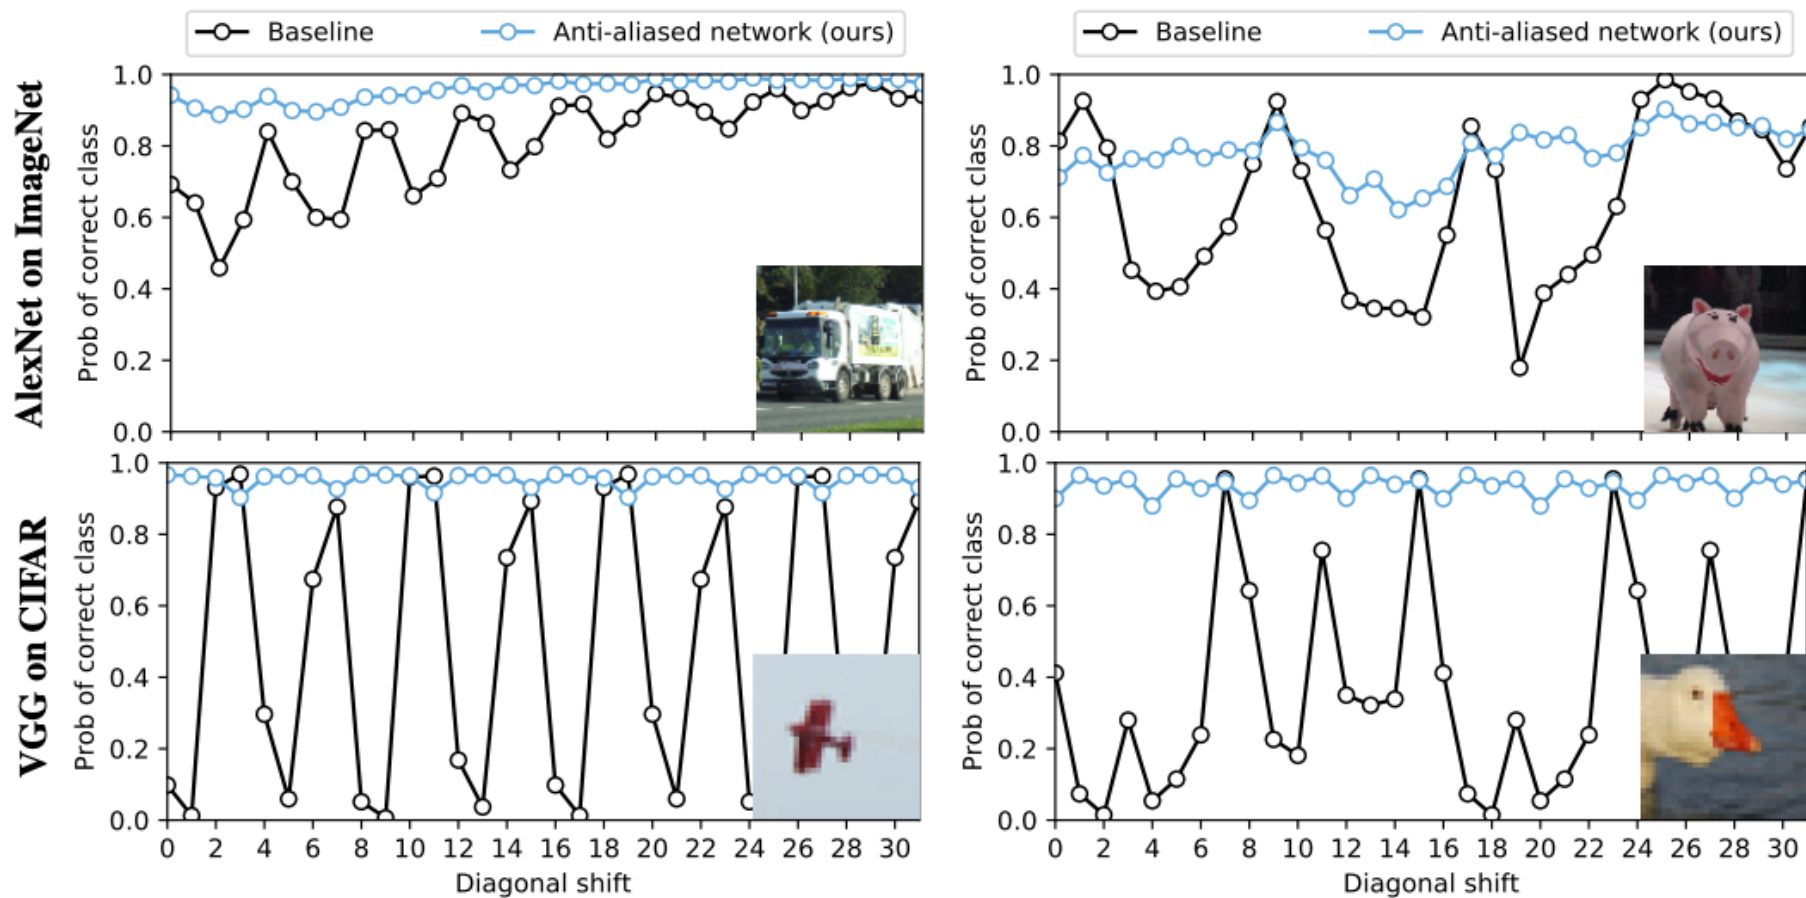



Figure 1. Classification stability for selected images. Predicted probability of the correct class changes when shifting the image. The baseline (black) exhibits chaotic behavior, which is stabilized by our method (blue). We find this behavior across networks and datasets. Here, we show selected examples using AlexNet on ImageNet (**top**) and VGG on CIFAR10 (**bottom**). Code and anti-aliased versions of popular networks are available at <https://richzhang.github.io/antialiased-cnns/>.


R. Zhang. [Making convolutional networks shift-invariant again](#). ICML 2019

Aliasing in neural networks



Nyquist-Shannon Sampling Theorem

When sampling a signal at discrete intervals, the sampling frequency must be at least **twice the maximum frequency of the input signal** to allow us to reconstruct the original perfectly from the sampled version



Anti-Aliasing

- How can we get rid of aliasing?
 - Sample more often (if you can)
 - Get rid of all frequencies that are greater than half the new sampling frequency
 - Will lose information, but that's better than aliasing
 - How to get rid of high frequencies?
 - Apply a **low-pass filter** (to be covered later)

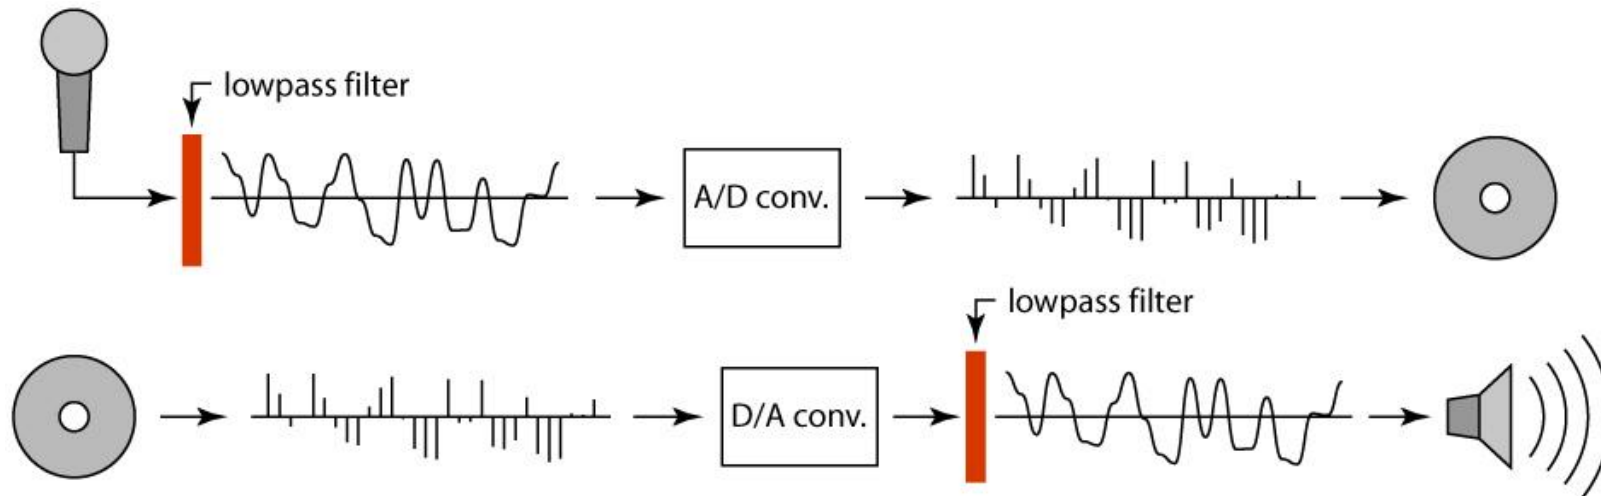



Image Resampling and Interpolation


Subsampling an Image

How do we reduce the size of an image by a factor of two?




Subsampling an image

How do we reduce the size of an image by a factor of two?



How about throwing away every other row and column to create a half-size image?

Subsampling without Pre-filtering



$1/2$




$1/4$ (2x zoom)



$1/8$ (4x zoom)

Subsampling with Pre-filtering



$1/2$



$1/4$




$1/8$

Image is smoothed with a *Gaussian filter* before subsampling

Subsampling with Pre-filtering


Image



Downsampled by 4




Smoothed then
downsampled by 4



Upsampling an Image


How do we *increase* the size of an image by a factor of two?



Let's increase the resolution of the sampling grid!


Upsampling an Image

How do we *increase* the size of an image by a factor of two?




Upsampling an Image

- How do we *increase* the size of an image by a factor of two?




Bilinear Interpolation

- Let $f(0, 0) = A, f(1, 0) = B, f(0, 1) = C, f(1, 1) = D$




Bilinear Interpolation

- Let $f(0, 0) = A, f(1, 0) = B, f(0, 1) = C, f(1, 1) = D$






Bilinear Interpolation


- Let $f(0, 0) = A, f(1, 0) = B, f(0, 1) = C, f(1, 1) = D$



Application: Demosaicing





Bilinear Interpolation More Generally



http://en.wikipedia.org/wiki/Bilinear_interpolation

Bilinear Interpolation More Generally



$$f(x, y) = w_{11}A + w_{21}B + w_{12}C + w_{22}D$$

$$w_{11} = \frac{(x_2 - x)(y_2 - y)}{(x_2 - x_1)(y_2 - y_1)}$$

$$w_{12} = \frac{(x_2 - x)(y - y_1)}{(x_2 - x_1)(y_2 - y_1)}$$


$$w_{21} = \frac{(x - x_1)(y_2 - y)}{(x_2 - x_1)(y_2 - y_1)}$$

$$w_{22} = \frac{(x - x_1)(y - y_1)}{(x_2 - x_1)(y_2 - y_1)}$$


http://en.wikipedia.org/wiki/Bilinear_interpolation

Bilinear interpolation: Basis function view


Interpolated function is sum of basis functions or “bumps” centered at the four adjacent grid points, weighted by the image values at the corresponding points




Bilinear basis function




Other kinds of interpolation




1D nearest-neighbour




Linear




Cubic



2D nearest-neighbour





Bilinear



Bicubic

Interpolation and function extrema

- When you use linear interpolation, extrema of the image function can only occur at the original sample points
- What about nonlinear interpolation?



Intensity Transformations


Intensity Transformations: Introduction

In intensity transformations, the output value of the pixel depends only on the input values of that pixel, not its neighbors.

Input image: $f(x)$


Output image: $g(x)$

Intensity transform $g(x) = T(f(x))$




E.g., identity transform $T(.) = 1$

```
import cv2 as cv
import numpy as np
f = cv.imread('images/emma_gray.jpg',
cv.IMREAD_GRAYSCALE)
t = np.arange(256, dtype=np.uint8)
g = t[f]
```



1. Explain the line $g = \text{cv.LUT}(f, t)$.
2. This can be done using NumPy only as $g = t[f]$. Explain this operation.
3. What is $T()$ here?

Example



$$g(x) = 255 - f(x)$$

Implementation:

```
t = np.arange(255, -1, -1, dtype=np.uint8)
```

Intensity Windowing


```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
```

```
c = np.array([(100, 50), (150, 200)])
```

```
t1 = np.linspace(0, c[0,1], c[0,0] + 1 - 0).astype('uint8')
t2 = np.linspace(c[0,1] + 1, c[1,1], c[1,0] - c[0,0]).astype('uint8')
t3 = np.linspace(c[1,1] + 1, 255, 255 - c[1,0]).astype('uint8')
```

```
transform = np.concatenate((t1, t2), axis=0).astype('uint8')
transform = np.concatenate((transform, t3), axis=0).astype('uint8')
print(len(transform))
```

```
img_orig = cv.imread('images/katrina.jpg', cv.IMREAD_GRAYSCALE)
image_transformed = cv.LUT(img_orig, transform)
```




Gamma Correction

$$g = f^\gamma, \quad f \in [0,1]$$


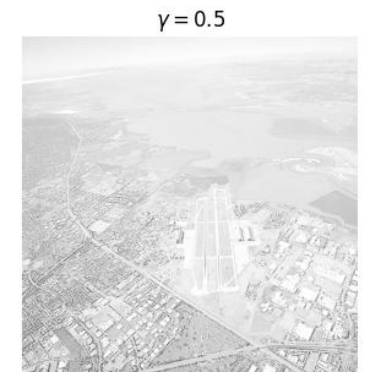

Values of γ such that $0 < \gamma < 1$ map a narrow range of dark pixels to a wider range of dark pixels.

$\gamma > 0$ has the opposite effect.

$\gamma = 1$ gives the identity transform.



```
t = np.array([(i/255.0)**(gamma)*255 for i in  
np.arange(0,256)]).astype(np.uint8)  
g = cv.LUT(f, t)
```



Histograms

1. We can represent the intensity distribution over the range of intensities $[0, 255]$, using a histogram.
2. If h is the histogram of a particular image, $h(r_k)$ gives us how many pixels have the intensity r_k .
3. The histogram of a digital image with gray values in the range $[0, L - 1]$ is a discrete function $h(r_k) = n_k$ where r_k is the k th gray level and n_k is the number of pixels having gray level r_k .
4. We can normalize the histogram by dividing by the total number of pixels n . Then we have an estimate of the probability of occurrence of level r_k , i.e., $p(r_k) = n_k/n$.
5. The histogram that we described above has L bins. We can construct a coarser histogram by selecting a smaller number of bins than L . Then several adjacent values of k will be counted for a bin.

Exercise

The figure shows a 3×4 image. The range of intensities that this image has is $[0, 7]$. Compute its histogram.

6	5	5	3
7	6	6	4
2	3	5	4

Image Properties through the Histogram

1. If the image is dark histogram will have many values in the left region, that correspond to dark pixels.
2. If the image is bright histogram will have many values in the right region, that correspond to bright pixels.
3. If a significant number of pixels are dark and a significant number of pixels are bright, the histogram will have two modes, one in the left region and the other in the right region.
4. A flat histogram signifies that the image has a uniform distribution of all intensities. Then, the contrast is high, and image will look vibrant.

Histogram Equalization

1. Photographers like to shoot pictures with a flat histogram, as such pictures are vibrant.
2. Histogram equalization is a gray-level transformation that results in an image with a more or less flat histogram.
3. We can take an image and make its histogram flat by using the operation called histogram equalization.

$$s = T(r) = (L - 1) \int_0^r p_r(w) dw$$

$$s = T(r) = \frac{L - 1}{MN} \sum_{j=0}^k n_j, \quad k = 0, \dots, L - 1$$

Example:

Suppose that a 3-bit image ($L = 8$) of size 64×64 pixels ($MN = 4096$) has the intensity distribution shown in Table 4, where the intensity levels are in the range $[0, L - 1] = [0, 7]$. Carry out histogram equalization.

r_k	n_k	$p_r(r_k) = n_k/MN$	$\sum_{j=0}^k n_j$	$\frac{(L-1)}{MN} \sum_{j=0}^k n_j$	Rounded
$r_0 = 0$	790	0.19			
$r_1 = 1$	1023	0.25			
$r_2 = 2$	850	0.21			
$r_3 = 3$	656	0.16			
$r_4 = 4$	329	0.08			
$r_5 = 5$	245	0.06			
$r_6 = 6$	122	0.03			
$r_7 = 7$	81	0.02			

r_k	n_k	$p_r(r_k) = n_k/MN$	$\sum_{j=0}^k n_j$	$\frac{(L-1)}{MN} \sum_{j=0}^k n_j$	Rounded
$r_0 = 0$	790	0.19	790		
$r_1 = 1$	1023	0.25			
$r_2 = 2$	850	0.21			
$r_3 = 3$	656	0.16			
$r_4 = 4$	329	0.08			
$r_5 = 5$	245	0.06			
$r_6 = 6$	122	0.03			
$r_7 = 7$	81	0.02			

r_k	n_k	$p_r(r_k) = n_k/MN$	$\sum_{j=0}^k n_j$	$\frac{(L-1)}{MN} \sum_{j=0}^k n_j$	Rounded
$r_0 = 0$	790	0.19	790		
$r_1 = 1$	1023	0.25	1813		
$r_2 = 2$	850	0.21	2663		
$r_3 = 3$	656	0.16	3319		
$r_4 = 4$	329	0.08	3648		
$r_5 = 5$	245	0.06	3893		
$r_6 = 6$	122	0.03	4015		
$r_7 = 7$	81	0.02	4096		

r_k	n_k	$p_r(r_k) = n_k/MN$	$\sum_{j=0}^k n_j$	$\frac{(L-1)}{MN} \sum_{j=0}^k n_j$	Rounded
$r_0 = 0$	790	0.19	790	1.350	1
$r_1 = 1$	1023	0.25	1813	3.098	3
$r_2 = 2$	850	0.21	2663	4.551	5
$r_3 = 3$	656	0.16	3319	5.672	6
$r_4 = 4$	329	0.08	3648	6.234	6
$r_5 = 5$	245	0.06	3893	6.653	7
$r_6 = 6$	122	0.03	4015	6.862	7
$r_7 = 7$	81	0.02	4096	7.000	7

Do the quiz.


Meeting at 2:30

Histogram Equalization Using OpenCV

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
```

```
f = cv.imread('images/shells.tif', cv.IMREAD_GRAYSCALE)
g = cv.equalizeHist(f)
```

Original Image



Histogram Equalization



Histogram Equalization Using the Formula

```
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt

f = cv.imread('images/shells.tif', cv.IMREAD_GRAYSCALE)

t = np.array([(L-1)/(M*N)*cdf[k] for k in range(256)],
dtype=np.uint8)
g = t[f]

g = cv.equalizeHist(f)
fig, ax = plt.subplots(1, 2, figsize=(12, 8))
ax[0].imshow(f, cmap='gray', vmin=0, vmax=255)
ax[0].set_title('Original Image')
ax[0].axis('off')
ax[1].imshow(g, cmap='gray', vmin=0, vmax=255)
ax[1].set_title('Histogram Equalization')
ax[1].axis('off')
plt.show()
```

Summary

1. Basics: digital image, color images, creating an image in Python, displaying images, increasing brightness, image planes
2. Intensity transformations:

$$g = t[f]$$

1. Identity, negative, intensity windowing
2. Gamma correction
3. Histograms, histogram equalization